

a Data

u_t : 19 executed buttons



4 training maps,
500 episodes each

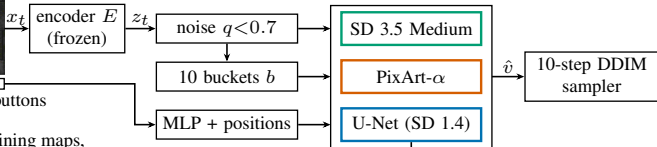


13 unseen arenas,
24 each

decoder D (stock)

b One recipe, three pretrained backbones

context and the noisy next latent stacked on channels

**c** Post-training

200k EMA weights

world model:
arena adaptation
(rank-16 LoRA,
8 episodes)

renderer:
decoder fine-tune
(MSE + 0.1 LPIPS)

d Evaluation

one tic ahead,
scene rows 0–207

