


a Data

u_t : 19 executed buttons



4 training maps,
500 episodes each



13 unseen maps,
24 each

decoder D (stock)

b Models and training

context and the noisy next latent stacked on channels

encoder E
(frozen)

z_t

noise $q < 0.7$

10 buckets b

MLP + positions

SD 3.5 Medium

PixArt- α

U-Net (SD 1.4)

$\hat{v}$

10-step DDIM
sampler

c Adaptation

200k EMA weights

world model: map adaptation
(rank-16 LoRA, 8 episodes)

decoder fine-tuning
(MSE + 0.1 LPIPS)

d Evaluation

one tic ahead, scene rows
0–207, against the raw
frame x_{t+1}

$\hat{z}_{t+1}$

decoder D

PSNR, LPIPS