


a Data x_t encoder E
(frozen) z_t  u_t : 19 executed buttons4 training maps,
500 episodes each13 unseen maps,
24 eachdecoder D (stock)**b Models and training**

context and the noisy next latent stacked on channels

noise $q < 0.7$ 10 buckets b

MLP + positions

SD 3.5 Medium

PixArt- α

U-Net (SD 1.4)

 $\hat{v}$ 10-step DDIM
sampler $\hat{z}_{t+1}$ decoder D

PSNR, LPIPS

directional score

c Adaptation

200k EMA weights

world model: map adaptation
(rank-16 LoRA, 8 episodes)decoder fine-tuning
(MSE + 0.1 LPIPS)**d Evaluation**one tic ahead, scene rows
0–207, against the raw
frame x_{t+1}