


**a Data** $x_t$ encoder  $E$   
(frozen) $z_t$  $u_t$ : 19 executed buttons4 training maps,  
500 episodes each13 unseen maps,  
24 eachdecoder  $D$  (stock)**b Models and training**

context and the noisy next latent stacked on channels

noise  $q < 0.7$ 10 buckets  $b$ 

MLP + positions

SD 3.5 Medium

PixArt- $\alpha$ 

U-Net (SD 1.4)

 $\hat{v}$ 10-step DDIM  
sampler**c Adaptation**

200k EMA weights

world model: map adaptation  
(rank-16 LoRA, 8 episodes)decoder fine-tuning  
(MSE + 0.1 LPIPS)**d Evaluation**one tic ahead, scene rows  
0–207, against the raw  
frame  $x_{t+1}$  $\hat{z}_{t+1}$ decoder  $D$ 

PSNR, LPIPS