

DoomShift: How Much Adaptation Do World Models Need Under Domain Shift?

Rohan Nagabhirava

Carnegie Mellon University
Pittsburgh, PA, United States
rnagabhi@andrew.cmu.edu

Keerthana Chirumamilla

Carnegie Mellon University
Pittsburgh, PA, United States
kchiruma@andrew.cmu.edu

Changliu Liu

Carnegie Mellon University
Pittsburgh, PA, United States
cliu6@andrew.cmu.edu

Abstract: World models have become capable simulators, but they are evaluated almost exclusively on held-out rollouts from the domain they were trained on. A small domain shift, such as a new map of the same game, a new texture set, a lighting change or a different cast of enemies, makes their outputs drift and lose realism, though the underlying physics and game engine are exactly the same. With DoomShift, a controlled version of this shift in Doom, we measure what it breaks and how much adaptation it takes to repair. We train three diffusion backbones, initialized from pretrained diffusion image models, on four maps of Doom and score each on in-domain data and on 13 other out-of-distribution maps that it never saw. All three keep their response to the controls and lose the scene: SD 3.5’s perceptual error (LPIPS) doubles, from 0.126 to a median 0.262, and its scene PSNR drops 3.3 dB, from 25.4 to a median 22.1 dB; the SD 1.4 U-Net, GameNGen’s backbone, and PixArt- α lose the same way. Very little adaptation repairs it: rank-16 adapters trained on 8 episodes of each unseen map with 2% of the original training updates recover 77% of SD 3.5’s lost perceptual quality and half of its lost PSNR, and 72% and 76% of the perceptual loss on the other two backbones.

Keywords: World models, Few-shot adaptation

1 Introduction

Game world models now render playable worlds from pixels and controls. GameNGen runs Doom inside a diffusion model [1], DIAMOND learns Atari and Counter-Strike in pixel space [2] and Matrix-Game 2 scales the recipe to open worlds [3]. Each of these models is evaluated only on held-out rollouts from the environment it was trained on: the reported score is a held-out trajectory drawn from its own training scenes and none has been shown to generalize to a new environment within the same game. This holds even for MultiGen, which trains a single Doom model across 100 procedurally generated maps with a pre-trained agent [4]: it reports PSNR, SSIM and LPIPS only on its training distribution, with no map held out for evaluation.

A learned simulator is useful only where it is faithful, and a robot or an agent moving to a new room, level or workspace needs to know how much of its world model still holds there and what it costs to make the rest hold. Game worlds let us pose this question with the engine held fixed and only the environment changed: a new map of the same game preserves the physics and the controls exactly, and only the layout, textures and lighting differ. Yet unchanged dynamics alone do not guarantee a faithful render: pretrained video and world models are typically moved to a new domain only through a further full fine-tune [5] or a dedicated adapter on the frozen backbone [6], and a model trained on a handful of maps tends to fall back on the appearance it already knows once the surroundings change. Pretraining one model on every map a deployment might meet is never complete, and training a separate model per map from scratch is prohibitively costly. We therefore explore what it takes to close the domain gap that opens when a world model meets a shift in which the underlying physics and game engine are exactly the same and only the surroundings, scenario or




Figure 1: On an unseen map the U-Net still moves and fires but renders the training maps’ masonry; eight adaptation episodes restore the map. Left: training maps (in distribution). Right: unseen map 7, zero-shot and after adaptation.

scene change: what the model retains, what it loses, and how few episodes and updates recover the rest.

In short, a world model moved to a new map keeps playing the game but draws the wrong world, and a small adapter trained on eight episodes draws the map back (Figure 1). We make three contributions:

- **DoomShift, a benchmark that isolates a scene shift.** Seventeen Doom maps run on exactly the same physics and game engine and differ only in their surroundings. We build the pipeline end to end (Figure 2): every-tic recordings, three diffusion world models initialized from pretrained diffusion image models and trained on four maps under one recipe, a decoder fine-tuned for the game, and a standard adapter for a new map, scored on the 13 maps the models never saw, zero-shot and after adaptation, against recorded frames and their decoder’s own reconstruction, which separates the model’s error from the renderer’s. Recordings, models, decoders and adapters are released (link withheld). This paper trains on 500 of the released episodes per training map and evaluates on 25 more.
- **A scene shift breaks appearance, not control.** On the unseen maps all three models keep their turn response (directional score 0.80 to 0.81, against 0.84 to 0.86 on the training maps) while their perceptual error roughly doubles (LPIPS 0.126 to a median 0.262 for SD 3.5) and their scene PSNR drops 2.8 to 3.3 dB. The loss is the same for a U-Net and two transformers from 0.6 to 2.3 B parameters, so across our three backbones it follows what the models were exposed to, not their architecture or scale. It is a step, not a slope: every unseen map is worse in LPIPS than every training map. And it lives in the model, since the decoder renders the unseen maps nearly as well as the training maps (Table 1, Figure 1).
- **How much adaptation it takes: very little.** A standard LoRA adapter on the frozen model (0.49 percent of the U-Net’s parameters), trained on eight episodes of the new map, recovers half of its lost perceptual quality in a median 250 updates. By 4k updates, 2 percent of pretraining and about one A4000 GPU-hour, it recovers 77 percent of SD 3.5’s lost perceptual quality and half of its lost PSNR; unchanged, the same procedure recovers 72 and 76 percent on the U-Net and PixArt- α . A full fine-tune with 205 times the trainable parameters recovers 87 percent against the adapter’s 71 on four maps (Tables 1 and 2, Figure 3).

2 Related work

Held-out evaluation of world models exists, but it holds out something other than a map of the same game: a new robot body [7], a new game or environment [6, 5] or a new building [8]. Doom maps have been held out, but only to test the agents that play them [9, 10], and game world models that do show new scenes generate them from an image, with no recorded frames to score against [11, 12]. To our knowledge, no published game world model has been scored against recorded frames on

maps of its own game that were deliberately held out from training. Our adapter is not new: a small adapter on a frozen video model moves it to a new setting in AVID [6] and adds action control in Vista [13], and AdaWorld adapts with a short full fine-tune [5], which we run as the comparator. None of them reports what adapting to one new scene costs in episodes and updates, which is what we measure.

3 Study design

Data. We collect the data for this research by running experiments with the Arnold agent [9] through ViZDoom [14] deathmatches: free-for-all matches where each player scores by killing the others. Arnold ships with 17 Doom maps; each has its own layout, textures and lighting. For every episode, we let Arnold play against 8 bots for 150 seconds. We record every tic, the engine’s 1/35 s simulation step, capturing a 320×240 frame and the 19 buttons the engine executed. We split the 17 maps into two groups. Four maps (2 to 5) are our training maps: we train on 500 episodes of each and score on 25 further validation episodes per map (512 windows). The other thirteen maps (1, 6 to 17) are never trained on: we record 24 new episodes on each, using 16 for adaptation and holding out 8 (256 windows) for scoring. In total we collect 2,412 episodes, about 96 hours of play and about 12.1 million frames.

Three backbones, one training recipe. We pick three pretrained backbones to ask three different questions. Our anchor is the SD 1.4 U-Net [15] (860M parameters): it is GameNGen’s backbone [1], so our in-distribution numbers come from the same network as the published Doom result. To test how different architectures perform at fixed scale, we swap in PixArt- α [16], a 628M transformer in the same latent space instead of a U-Net; architecture turns out not to change what survives the shift: in distribution their scene PSNR differs by 0.02 dB and their LPIPS by 0.001, and they lose the same amount on the unseen maps (Table 1). To test how performance changes with scale, we use SD 3.5 Medium [17], a widely adopted open MMDDiT (2.27B parameters, 16-channel latents). Its 16-channel latents give its decoder a higher reconstruction score (31.8 against 28.6 dB on the training maps).

We turn all three into next-tic world models under one shared recipe (Figure 2). We give each model 32 context tics so it can see motion, train with v -prediction [18] for stable training across noise levels and add context noise augmentation as in GameNGen [1] so the model tolerates its own rollout errors. We train each backbone for 200k updates at batch size 32, a step count we match across backbones, keep an EMA at 0.9999 in fp32 (a bf16 EMA underflows and freezes) and sample with 10-step DDIM [19], which we found to be a strong balance between quality and efficiency (sampler sweep in the supplementary material).

Metrics. We score every predicted frame against four questions: how accurate is it against the raw ground-truth frame, how realistic does it look and does it turn the right way when we mirror the controls? We score the part of each frame that shows the world (rows 0 to 207) and leave out the HUD rows below it, which do not change with the map.

For accuracy, we compute PSNR against the raw frame, and for realism, LPIPS [20]. Absolute PSNR alone does not compare maps, because it rises with how static a recording is: copying the last context frame forward, the persistence baseline of video prediction [21], already scores 19.8 dB on the unseen maps (18.5 to 22.5 per map). A map’s *deficit* is its score’s distance from the same model’s in-distribution level on the training maps; its *excess gap* is that distance in PSNR. As a diagnostic, the decoder applied to the ground-truth latent scores 27.3 dB on the unseen maps and 28.6 on the training maps, so the decoder itself loses little on the shift. For turning, we compute the *directional* score: the fraction of turning windows whose predicted view turns the other way once we swap left and right in the newest control.

Decoder. Every backbone is trained on latents from its frozen encoder, so the latent space is the contract between encoder, world model and decoder: changing the encoder would move every training target and invalidate the trained backbones, the adapters and the decoder’s reconstruction scores.




Figure 2: What we built: (a) the benchmark, (b) three backbones under one recipe, (c) adaptation, (d) evaluation against the raw ground-truth frame.

The decoder only renders: it sits after the model, and its own error is what its reconstruction score measures, so fine-tuning it changes what we can see, not what the model knows. We fine-tune the SD 1 decoder D on training-map frames paired with their frozen-encoder latents, using MSE plus 0.1 LPIPS for 3,486 updates and gating it on held-out training-map frames (the supplementary material); it never sees an unseen map. On 200 validation frames it raises the scene reconstruction from 26.9 to 28.4 dB and lowers its LPIPS from 0.092 to 0.060, and it reconstructs the unseen maps nearly as well as the training maps (median 27.3 against 28.6 dB) while the model’s prediction loses 2.9 dB, so the shift lives in the model, not the renderer; the LPIPS term is there because MSE alone more than triples scene LPIPS (0.092 to 0.339). We use it for every pixel score of the U-Net and PixArt- α unless stated otherwise; SD 3.5 uses its own decoder, fine-tuned the same way.

Adaptation as a measurement. To measure how much adaptation a map needs, we keep the trained model frozen and train a standard adapter on a few episodes of that map, then score the adapted model on held-out episodes at each update budget. For each unseen map we adapt each backbone’s 200k EMA weights with rank-16 LoRA [22] on every attention projection of the frozen backbone [13] and train the control MLP, input projection and noise-bucket embedding in full (4.2M parameters, 0.49 percent of the U-Net; 25.9M, 4.1 percent, of PixArt- α ; 40.6M, 1.8 percent, of SD 3.5) [23]. We train on 8 adaptation episodes (for the U-Net about 1 A4000 GPU-hour per map for 4k updates) and score the non-EMA weights at 0, 50, 100, 150, 250, 500, 1k, 2k, 4k and 8k updates (PixArt- α to 4k; SD 3.5 at 0, 250, 500, 1k, 2k and 4k), with the control checks (forgetting, the scene PSNR on the training maps, and the directional check) at 0 and 8k. 4k updates is the elbow of the curve, and we report it everywhere: across the 13 unseen maps it captures most of the benefit (on average 96 percent of each map’s 8k gain; Figure 3) at the smallest budget, which is why adaptation costs 2 percent of the original training updates; 8k serves as the check. As a comparator we fine-tune all 860.5M parameters of the U-Net at 2×10^{-5} on maps 6, 7, 8 and 16 with the same eight episodes and score it at 4k; it runs on an RTX A6000 (49 GB) at 1.24 updates per second, since it does not fit the 16 GB A4000 the adapter uses. A map’s budget is the first grid point at which it recovers at least half of its zero-shot LPIPS rise over the training maps; maps that do not by 8k are right-censored [24].

4 Results

We report what a model loses zero-shot, what adaptation recovers and with what budget, and what predicts how far a map gets.

The shift breaks the scene, not the controls. What does a model keep and what does it lose when it meets a map it never saw? On unseen maps the turn response survives (U-Net directional scores 0.73 to 0.86 per map against 0.79 to 0.91 in distribution; Table 1). The appearance does not: the U-Net’s scene PSNR falls from 25.2 dB on the training maps to a median 22.3 dB (20.5 to 24.7 per map) and its LPIPS rises from 0.158 to 0.303 (0.229 to 0.401). Pooled over the 13 unseen maps,

Table 1: Scene PSNR and LPIPS one tic ahead against the raw frame and the directional score: training maps (pooled over 512 validation windows) and unseen maps (medians over the 13 unseen maps), each backbone through its own fine-tuned decoder; directional pooled over the turning windows (ground-truth frames 0.885, 0.892).

Model	Maps	PSNR (dB)	LPIPS	Directional
SD 1.4 U-Net	training	25.20	0.158	0.855
	unseen	22.30	0.303	0.804
PixArt- α	training	25.18	0.159	0.842
	unseen	22.38	0.285	0.808
SD 3.5 Medium	training	25.38	0.126	0.840
	unseen	22.09	0.262	0.800




Figure 3: The shift and its repair, per unseen map and per backbone. (a) Scene PSNR and (b) LPIPS zero-shot (open) to 4k adapter updates (filled) for the U-Net, PixArt- α and SD 3.5, with each backbone’s in-distribution level on the training maps (dashed); (c, d) medians over the 13 unseen maps against updates.

one predicted frame in five has an LPIPS above 0.4, against one in 73 on the training maps; such frames concentrate in maps 16, 17, 13 and 1 (32 to 59 percent of their windows), while seven maps stay below 10 percent. The decoder is not the limit: it reconstructs the unseen maps nearly as well as the training maps (median 27.3 against 28.6 dB, LPIPS 0.05 to 0.09). The three backbones lose the same things (Table 1; the supplementary material): scene PSNR drops 2.9, 2.8 and 3.3 dB and LPIPS rises by 0.145, 0.126 and 0.136, so across our three backbones the shift looks like a property of the task, not of the architecture or the scale. It is a step, not a slope: every unseen map has a worse LPIPS than every training map for all three backbones (the supplementary material). What the model keeps is how actions act on the world: zero-shot on map 7 the weapon fires and the view advances as in the ground truth (Figure 1) and the turn response holds; the loss sits in the scene, whose rows fall 2.9 dB (25.2 to a median 22.3) while the map-independent HUD rows fall 0.2 dB (29.7 to 29.5).

How much adaptation: updates, episodes, parameters. We next adapt a small subset of each backbone’s parameters on each unseen map and ask what it takes to recover what was lost. Eight episodes and 4k updates raise the median scene PSNR from 22.3 to 23.6 dB and lower the median LPIPS from 0.303 to 0.210 (Figure 3, Table 2); the 8k check adds little (23.7 dB, 0.204). A map’s budget, the first of the grid reads (50, 100, 150, 250, 500, 1k, 2k, 4k and 8k updates) at which it

Table 2: What closing the gap takes. The U-Net’s LoRA by group of unseen maps, terciles of their zero-shot LPIPS rise over the training maps (hard: the largest rise), and over all 13 unseen maps; then the U-Net’s full fine-tune on the four comparator maps 6, 7, 8 and 16 (against the adapter’s 22.88 dB / 0.209 and 71 percent on the same maps), and the same LoRA procedure on PixArt- α and SD 3.5 over all 13 unseen maps. Medians over each row’s maps; non-EMA adapter weights; scene rows against the raw frame, each backbone through its fine-tuned decoder. Recovered: the share of the zero-shot LPIPS rise that the adapter undoes and of the lost scene PSNR that it regains, shares per map, then the median. —: not scored (the PixArt- α and SD 3.5 grids stop at 4k). Budget: the first grid point at which a map recovers at least half of its zero-shot LPIPS rise ($>4k$: censored at the grid’s last read; the full fine-tune is scored at 0 and 4k only, so $\leq 4k$; nine of SD 3.5’s maps cross at 250, the first point of its grid; where a group’s two middle budgets differ, both are given).

Unseen maps (by zero-shot LPIPS rise)	Scene PSNR (dB) $\uparrow$			Scene LPIPS $\downarrow$			Recovered (%)		Budget to half the rise
	zero-shot	4k	8k	zero-shot	4k	8k	LPIPS	PSNR	
Hard (1, 13, 16, 17)	21.65	22.56	22.60	0.391	0.233	0.226	65	29	250–2k
Medium (6, 8, 12, 14, 15)	22.80	23.98	24.04	0.303	0.210	0.204	68	48	250
Easy (7, 9, 10, 11)	23.38	24.92	24.99	0.255	0.170	0.165	87	173	150
All 13	22.30	23.60	23.69	0.303	0.210	0.204	72	48	250
U-Net full fine-tune (6, 7, 8, 16)	21.71	23.11	—	0.292	0.182	—	87	46	$\leq 4k$
PixArt- α LoRA (all 13)	22.38	23.72	—	0.285	0.205	—	76	48	250; 1 of 13 censored
SD 3.5 LoRA (all 13)	22.09	24.00	—	0.262	0.165	—	77	50	250

recovers half of its zero-shot rise in LPIPS over the training maps, has a median of 250 updates (100 to 8k per map, none censored; 12 of the 13 maps cross by 4k), and half of each map’s 8k gain arrives within 100 updates on average. By the same 4k updates it recovers a median 72 percent of each map’s rise in LPIPS over the training maps (at least 48 percent) and a median 48 percent of its lost scene PSNR (over 100 percent on maps 8, 9 and 10, whose static recordings already sit near the training maps’ PSNR). About a decibel of the PSNR still missing against the training maps is the decoder’s own reconstruction drop on the unseen maps (27.3 against 28.6 dB), which no adapter of the world model can recover. The procedure carries over unchanged to the other backbones. On PixArt- α the adapter raises the median scene PSNR on the 13 unseen maps from 22.4 to 23.7 dB and lowers LPIPS from 0.285 to 0.205 by 4k: it recovers a median 76 percent of the rise in LPIPS and 48 percent of the lost PSNR with a median budget of 250 updates (map 1 censored at its 4k grid), against 72 and 48 percent and 250 updates for the U-Net. On SD 3.5 the same adapter raises the median scene PSNR on the 13 unseen maps from 22.1 to 24.0 dB and lowers LPIPS from 0.262 to 0.165 by 4k: it recovers a median 77 percent of the rise in LPIPS over its in-distribution 0.126 (at least 54 percent) and 50 percent of the lost PSNR. Its median budget is also 250 updates, but nine of its 13 maps cross at 250, the first point of its grid, so the rule cannot resolve finer there. A full fine-tune of all 860M parameters on the same eight episodes does better on each of the four comparator maps 6, 7, 8 and 16: at 4k it reaches a median scene PSNR of 23.1 dB and LPIPS of 0.182 against the adapter’s 22.9 dB and 0.209 (0.20 to 0.29 dB and 0.017 to 0.030 in LPIPS per map) and recovers a median 87 percent of the rise in LPIPS against the adapter’s 71. It trains 205 times as many parameters and needs a 49 GB RTX A6000, on which 4k updates take a median 0.8 hours, where the adapter trains on a 16 GB A4000 in 1.1 hours. Data matters less than the first updates: on the four maps of the data ladder, one episode already gives most of the gain of sixteen (the supplementary material). The adapters trade some in-distribution skill for it: the training maps’ scene PSNR against the decoded ground truth falls on all 13 unseen maps at the 8k check (median 0.59 dB; stock decoder), while the directional score rises from 0.81 to 0.84.

5 Conclusion

The claims cover 13 unseen maps of one game, one agent, one-tic scoring and mostly one seed. One-tic quality does not guarantee stable rollouts (the supplementary material). Next are other Doom maps and other games, to see whether the same procedure transfers; detecting a shift from the model’s own zero-shot scores so that it repairs itself; closed-loop evaluation of the adapted models;

and agents trained inside the adapted model. The pattern we measure, control kept, appearance lost, most of it relearned from a few episodes, is what a deployed world model would need to detect and repair on its own, and DoomShift is a testbed for that loop.

Acknowledgments

[tbd]

References

- [1] D. Valevski, Y. Leviathan, M. Arar, and S. Fruchter. Diffusion models are real-time game engines. In *International Conference on Learning Representations (ICLR)*, 2025. URL <https://arxiv.org/abs/2408.14837>.
- [2] E. Alonso, A. Jelley, V. Micheli, A. Kanervisto, A. Storkey, T. Pearce, and F. Fleuret. Diffusion for world modeling: Visual details matter in Atari. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. URL <https://arxiv.org/abs/2405.12399>.
- [3] X. He, C. Peng, Z. Liu, B. Wang, Y. Zhang, Q. Cui, F. Kang, B. Jiang, M. An, Y. Ren, B. Xu, H.-X. Guo, K. Gong, S. Wu, W. Li, X. Song, Y. Liu, Y. Li, and Y. Zhou. Matrix-Game 2.0: An open-source, real-time, and streaming interactive world model. *arXiv preprint arXiv:2508.13009*, 2025.
- [4] R. Po, D. J. Zhang, A. Hertz, G. Wetzstein, N. Wadhwa, and N. Ruiz. MultiGen: Level-design for editable multiplayer worlds in diffusion game engines. *arXiv preprint arXiv:2603.06679*, 2026.
- [5] S. Gao, S. Zhou, Y. Du, J. Zhang, and C. Gan. AdaWorld: Learning adaptable world models with latent actions. In *International Conference on Machine Learning (ICML)*, 2025. URL <https://proceedings.mlr.press/v267/gao25u.html>.
- [6] M. Rigter, T. Gupta, A. Hilmkil, and C. Ma. AVID: Adapting video diffusion models to world models. *Reinforcement Learning Journal*, 6:737–764, 2025. URL <https://rlj.cs.umass.edu/2025/papers/Paper64.html>.
- [7] Y. Chen, J. Yang, Y. Xu, Q. Ma, K. He, P. Li, K. Wang, Z. He, X. Wu, J. Liu, N. Liu, Y. Huang, and L. Wang. XEWorld: Can action-conditioned world models generalize to unseen robot embodiments? *arXiv preprint arXiv:2608.05799*, 2026.
- [8] J. Y. Koh, H. Lee, Y. Yang, J. Baldridge, and P. Anderson. Pathdreamer: A world model for indoor navigation. In *IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021. doi:10.1109/ICCV48922.2021.01447.
- [9] G. Lample and D. S. Chaplot. Playing FPS games with deep reinforcement learning. In *AAAI Conference on Artificial Intelligence*, 2017. doi:10.1609/aaai.v31i1.10827.
- [10] M. Wydmuch, M. Kempka, and W. Jaśkowski. ViZDoom competitions: Playing Doom from pixels. *IEEE Transactions on Games*, 11(3):248–259, 2019. doi:10.1109/TG.2018.2877047.
- [11] Z. Tong, Y. Jin, H. Lai, Z. Wang, Z. Xing, K. Cheng, H. Xu, Z. Pu, S. Zhu, R. Feng, J. Zhao, Y. Zhang, H. Tang, and L. Shao. SCOPE: Simulating cross-game operations in playable environments for FPS world models. *arXiv preprint arXiv:2605.23345*, 2026.
- [12] J. Bruce, M. Dennis, A. Edwards, J. Parker-Holder, Y. Shi, E. Hughes, M. Lai, A. Mavalankar, R. Steigerwald, C. Apps, Y. Aytar, S. Bechtle, F. Behbahani, S. Chan, N. Heess, L. Gonzalez, S. Osindero, S. Ozair, S. Reed, J. Zhang, K. Zolna, J. Clune, N. de Freitas, S. Singh, and T. Rocktäschel. Genie: Generative interactive environments. In *International Conference on Machine Learning (ICML)*, 2024. URL <https://proceedings.mlr.press/v235/bruce24a.html>.
- [13] S. Gao, J. Yang, L. Chen, K. Chitta, Y. Qiu, A. Geiger, J. Zhang, and H. Li. Vista: A generalizable driving world model with high fidelity and versatile controllability. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024. URL <https://arxiv.org/abs/2405.17398>.

- [14] M. Kempka, M. Wydmuch, G. Runc, J. Toczec, and W. Jaśkowski. ViZDoom: A Doom-based AI research platform for visual reinforcement learning. In *IEEE Conference on Computational Intelligence and Games (CIG)*, 2016. doi:10.1109/CIG.2016.7860433.
- [15] R. Rombach, A. Blattmann, D. Lorenz, P. Esser, and B. Ommer. High-resolution image synthesis with latent diffusion models. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022. doi:10.1109/CVPR52688.2022.01042.
- [16] J. Chen, J. Yu, C. Ge, L. Yao, E. Xie, Y. Wu, Z. Wang, J. Kwok, P. Luo, H. Lu, and Z. Li. PixArt- α : Fast training of diffusion transformer for photorealistic text-to-image synthesis. In *International Conference on Learning Representations (ICLR)*, 2024. URL <https://arxiv.org/abs/2310.00426>.
- [17] P. Esser, S. Kulal, A. Blattmann, R. Entezari, J. Müller, H. Saini, Y. Levi, D. Lorenz, A. Sauer, F. Boesel, D. Podell, T. Dockhorn, Z. English, and R. Rombach. Scaling rectified flow transformers for high-resolution image synthesis. In *International Conference on Machine Learning (ICML)*, 2024. URL <https://proceedings.mlr.press/v235/esser24a.html>.
- [18] T. Salimans and J. Ho. Progressive distillation for fast sampling of diffusion models. In *International Conference on Learning Representations (ICLR)*, 2022. URL <https://arxiv.org/abs/2202.00512>.
- [19] J. Song, C. Meng, and S. Ermon. Denoising diffusion implicit models. In *International Conference on Learning Representations (ICLR)*, 2021. URL <https://arxiv.org/abs/2010.02502>.
- [20] R. Zhang, P. Isola, A. A. Efros, E. Shechtman, and O. Wang. The unreasonable effectiveness of deep features as a perceptual metric. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2018. doi:10.1109/CVPR.2018.00068.
- [21] M. Mathieu, C. Couprie, and Y. LeCun. Deep multi-scale video prediction beyond mean square error. In *International Conference on Learning Representations (ICLR)*, 2016. URL <https://arxiv.org/abs/1511.05440>.
- [22] E. J. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, and W. Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations (ICLR)*, 2022. URL <https://arxiv.org/abs/2106.09685>.
- [23] E. Xie, L. Yao, H. Shi, Z. Liu, D. Zhou, Z. Liu, J. Li, and Z. Li. DiffFit: Unlocking transferability of large diffusion models via simple parameter-efficient fine-tuning. In *IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023. doi:10.1109/ICCV51070.2023.00390.
- [24] M. E. Taylor and P. Stone. Transfer learning for reinforcement learning domains: A survey. *Journal of Machine Learning Research*, 10(56):1633–1685, 2009. URL <https://www.jmlr.org/papers/v10/taylor09a.html>.